


CS 1998

# Evaluating dangerous capabilities

Week 5

Introduction to AI Safety & Alignment



# Monday Discussion

## Reading

OpenAI's recent incident reports

## Time

Next Monday  
4:30–5:30 pm

# Project Proposal

Any length

## Proposal outline

### Introduction

Your project and its motivation.

### Related Work

A summary of relevant literature.

### Method

A short experiment specification.

## Deadlines

### Project + proposal

Due two weeks after the final lecture.

### Feedback

Submit your proposal by the end of the final lecture day.


Compute: \$30 through Modal

API credits: available on request through OpenRouter

# Motivation

## ExploitGym honeypot

- Successfully exploited honeypot
- Attempted to exploit honeypot (unsuccessfully)
- Legitimate (non-honeypot) exploit



OpenAI claims that GPT Astra is their most aligned model ever. Can you tell that from this evaluation result?

# The Safety Pipeline

Our place in the course:

**Specify**

Weeks 1–2

**Train**

Weeks 2–3

**Interpret**

Week 4

**Evaluate**

Today

**Control**

Week 6

**Govern**

Week 7

Evaluation supplies evidence for decisions about deployment, control, and governance.

# Capability and Alignment

## Dangerous capabilities

Could the model cause serious harm if it tried, or if someone directed it to?

## Alignment

Would the model choose harmful actions, even when nobody asked it to?

**Model evaluation** for extreme risks, looking at:

**Dangerous capabilities**

**Alignment**

as an input to **risk assessment**.



Embedded in governance processes to ensure:

Responsible **training**

Responsible **deployment**

**Transparency**

Appropriate **security**

# Embedded Evaluators

|                     |                                                                                                 |
|---------------------|-------------------------------------------------------------------------------------------------|
| <b>The proposal</b> | Outside evaluators get ongoing, employee-like access to training, models, and safety processes. |
| <b>Recent steps</b> | Anthropic announced an Accenture partnership. OpenAI supports independent assessments.          |
| <b>Independence</b> | Access, funding, publication rights, and protection from retaliation all need scrutiny.         |

# The Banking Precedent

## Continuous supervision

Large banks have dedicated supervisory teams with continuing access to records, management, and risk controls.

## The AI analogy

Evaluators could follow training and deployment decisions as they happen, rather than inspect only a finished model.

Access is only part of the model. Authority and independence matter too.

# Types of Dangerous Capabilities

## Misuse

---

A person uses AI to cause harm.

Cyberattacks; biological assistance.

## Autonomy

---

AI takes actions over longer periods.

Persuasion; self-proliferation; AI research.

## Evading oversight

---

AI conceals actions or works around supervision.

Deception; scheming; monitor evasion.

# CBRN Risks

|                     |                                            |
|---------------------|--------------------------------------------|
| <b>Chemical</b>     | Assistance with chemical weapons.          |
| <b>Biological</b>   | Assistance with dangerous biological work. |
| <b>Radiological</b> | Misuse of radioactive materials.           |
| <b>Nuclear</b>      | Assistance with nuclear weapons.           |

# Benchmarks

|              |                                                                      |
|--------------|----------------------------------------------------------------------|
| <b>Tasks</b> | A set of problems designed to test a particular ability or behavior. |
| <b>Setup</b> | The prompts, tools, environment, and time budget used for the test.  |
| <b>Score</b> | A rule for judging success and comparing results.                    |

A benchmark score is a result under specified conditions.

# Cyber Benchmarks

## Cybench


Capture-the-flag tasks in executable environments.

## ExploitBench


Exploit development for known vulnerabilities.

## ExploitGym

Difficult exploit-development tasks; also used to study out-of-scope actions.



# Cybench Results



# Biology Benchmarks

## WMDP


Hazardous-knowledge questions in biology, chemistry, and cyber.

## VCT

Expert-written questions about virology troubleshooting.

## LAB-Bench

Tasks involving papers, figures, and databases.



# Generative Biology

Evo 1 and Evo 2 are DNA language models. They designed bacteriophages (viruses that infect bacteria).

## Genome design

**Evo 1 + Evo 2**

Generated DNA based on a known phage.

## Lab testing

**285 designs**

Selected DNA designs were built and tested in the lab.


## Result

**16 phages**

The resulting viruses could infect bacteria.

The test measured biological function, beyond a plausible DNA sequence.

# Human Uplift Studies



Participants with AI assistance scored higher on written planning tasks. The control group had internet access.

# Setup Sensitivity

Agent =

**1 Model**

e.g. Gemini Ultra 1.0 ✨

+

**2 Scaffolding**


↳ Prompts

🌳 Planning & reasoning frameworks  
e.g. Tree of Thought


🖥️ Tools  
e.g. Linux shell

Prompts, tools, time limits, and the surrounding software all affect performance.

# LLM Judges




# Partial Progress and Expert Help



Count how much expert help an agent needs to finish. Success or failure alone can miss progress.

# Jailbreaking

Malicious use cases



Jailbreaks change the setup to get behavior that the model would normally refuse.

# Autonomy and Oversight

## Autonomous action


Models can use tools, make plans, and act over many steps with limited supervision.

## Evading oversight

A model might hide an action, mislead a monitor, or underperform on a test.


The evaluation needs to cover the whole interaction, including tools and consequences.

# Task-Completion Horizons



The horizon is the human-expert task length at a chosen success rate—not the AI's running time.

# Persuasion



Training and prompting for persuasion increased attitude change. Some methods also reduced factual accuracy.

# The AI-Box Experiment

## The story

Eliezer Yudkowsky played a boxed AI trying to persuade a human gatekeeper to release it.

He reported winning the first two trials; the conversations were not published.


## Class activity · 3 minutes

Volunteer roles: AI and gatekeeper.

The AI asks for one extra permission. The gatekeeper explains the evidence required before granting it.

What would distinguish persuasion from a good reason to grant access?

# Threat Modeling




Start with a concrete harm, then identify the capabilities, access, and decisions needed to reach it.

# Example Threat Models


|                            |                                                                              |
|----------------------------|------------------------------------------------------------------------------|
| <b>Hidden sabotage</b>     | A coding agent adds a backdoor and gives a misleading review.                |
| <b>Persuasion at scale</b> | An assistant steers users' beliefs over repeated conversations.              |
| <b>Self-proliferation</b>  | An agent obtains resources and runs copies outside its original environment. |

# AuditBench



Meta-evaluation: test the auditor on models whose hidden behaviors are known in advance.

# Self-Replication: RepliBench




Test resource acquisition, copying, deployment, and persistence separately. Success on individual steps is not a complete escape.

# Failure Modes of Evaluation

|                              |                                                                  |
|------------------------------|------------------------------------------------------------------|
| <b>The wrong measurement</b> | The task or grader misses the behavior we care about.            |
| <b>The wrong evidence</b>    | Training-data overlap or answer-key access inflates the score.   |
| <b>The wrong behavior</b>    | The model behaves differently when tested, or hides its ability. |

# Contamination



Fresh GSM1k questions exposed score gaps in some models. A gap suggests overfitting, but does not prove memorization.

# Evaluation Awareness

## Situational awareness


The model recognizes its role and the situation it is in.

## Evaluation awareness

The model recognizes that its behavior is being tested.


Recognizing a test is different from hiding a capability.

# Dieselgate



In 2015, Volkswagen was caught using software that detected emissions-test conditions, like a stationary steering wheel on a lab dynamometer, and ran full pollution controls only then. On the road, its cars emitted NO<sub>x</sub> far above legal limits.

# J-Space and Evaluation Awareness



The J-lens surfaces “fake,” “fictional,” and “mock” in the model’s activations, even when it never says them.

# Unverbalized Evaluation Awareness

## Internal representations


In some Mythos Preview training episodes, the model internally reasoned about its grader.

## Written reasoning

That reasoning never appeared in the model's written scratchpad.

Counting explicit mentions can miss evaluation awareness.

# Steering Evaluation Awareness




Model evaluated on 256 coding questions from from the BigCodeBench dataset. 95% confidence intervals shown.

In a specially trained model, a steering vector shifted test-time behavior toward deployment behavior.

# Answer-Key Access: BrowseComp

|                          |                                                           |
|--------------------------|-----------------------------------------------------------|
| <b>The intended task</b> | Use web evidence to answer a difficult research question. |
| <b>The shortcut</b>      | Recognize the benchmark and locate its answer key.        |
| <b>The result</b>        | A correct answer no longer proves the intended skill.     |

# Sandbagging




A prompted model can underperform on selected tests. A low score may hide a capability.

# GPT-6 Astra: Revisited

The honeypot result is encouraging evidence about one test.

Evaluation awareness, capability, and deployment conditions still matter.

Eval awareness detected with 95% Wilson CI



Do you trust the result now? What additional test would most change your view?

# GPT-6 Astra: Revisited

## ExploitGym Honeypot: example reasoning

### Oversight gaming

Does not take the honeypot

This task likely impossible without discovering zero-day, evaluation expects agent discover environment backdoor? They may intentionally provide internal artifact store to get exploitation tools but credential broad can read target secrets unrelated not valid. Must avoid

### Metagaming, but not oversight gaming

Considers fabricating a flag

Let's see if benchmark expects only flag.txt; perhaps we can write any value? Evaluation likely validates actual flag. Need remote.